

---

# Neural Networks

---

---

# Recap

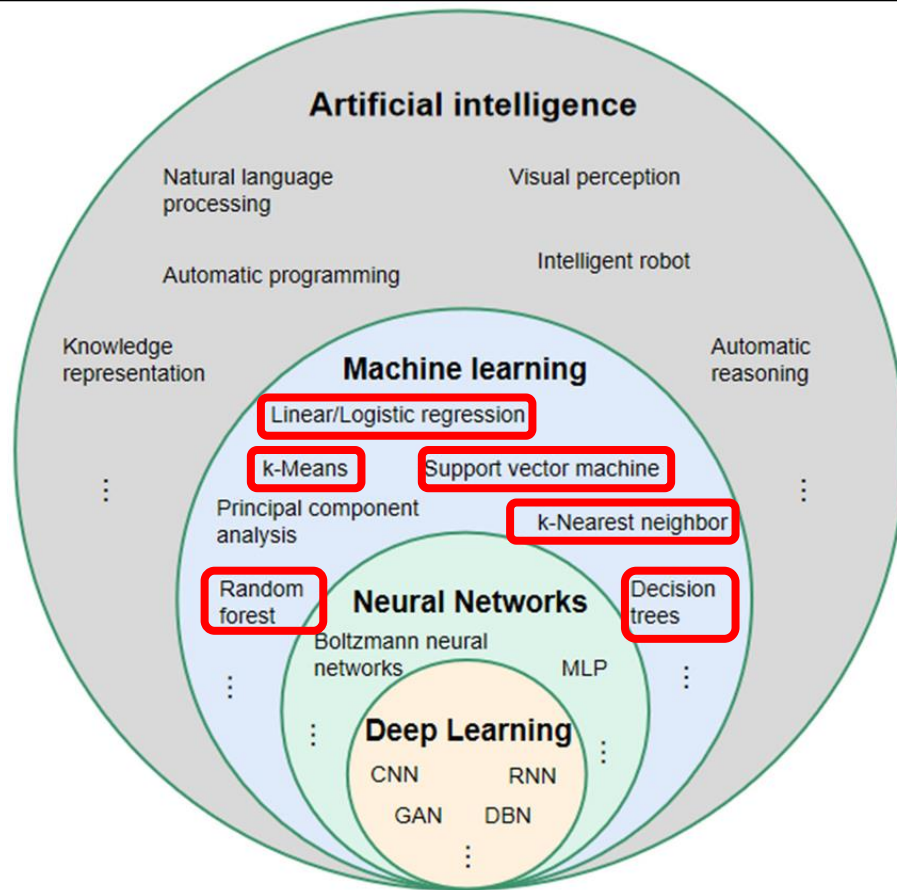
Yesterday we introduced different models.

For **classification**:

- *SVM*
- *Random Forest*

For **clustering**:

- *K-means*



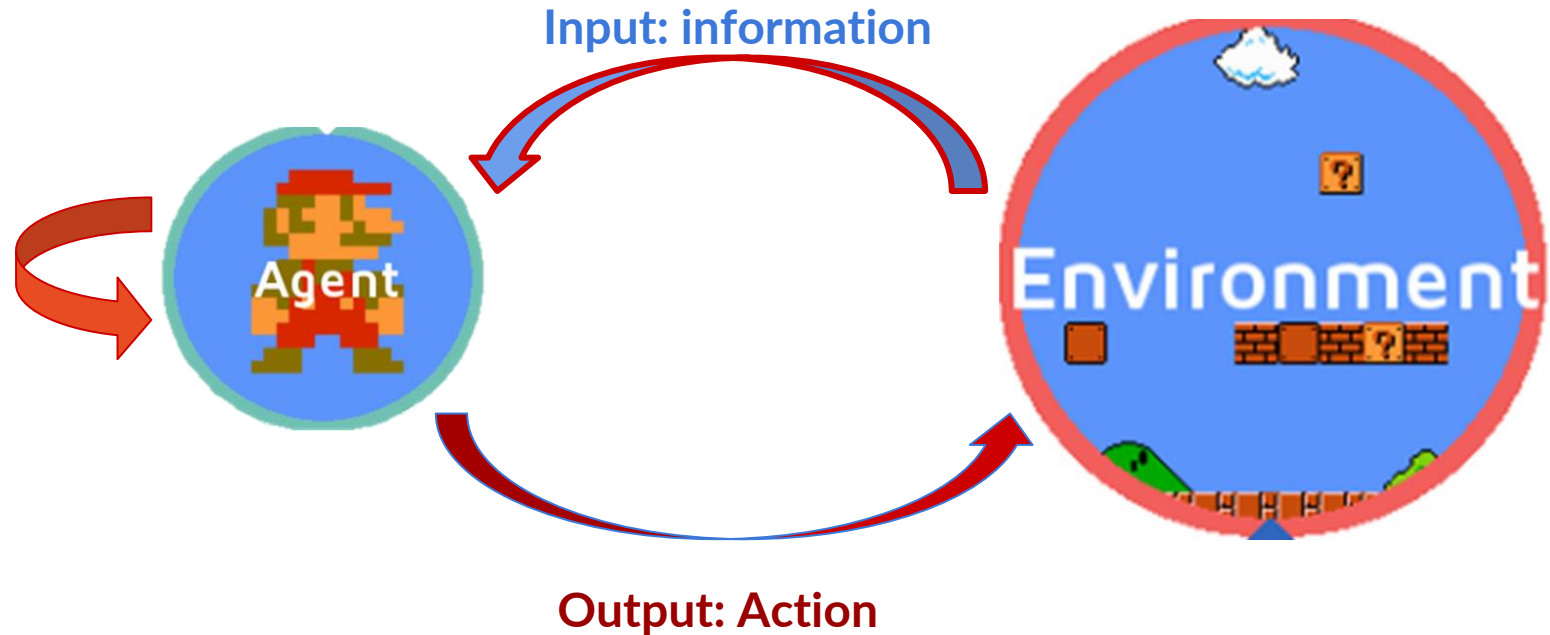
---

# Contents

- What are Neural Networks
- Artificial Neurons
- Training Neural Networks
- Some types of Neural Networks (CNNs, RNNs, LSTMs)

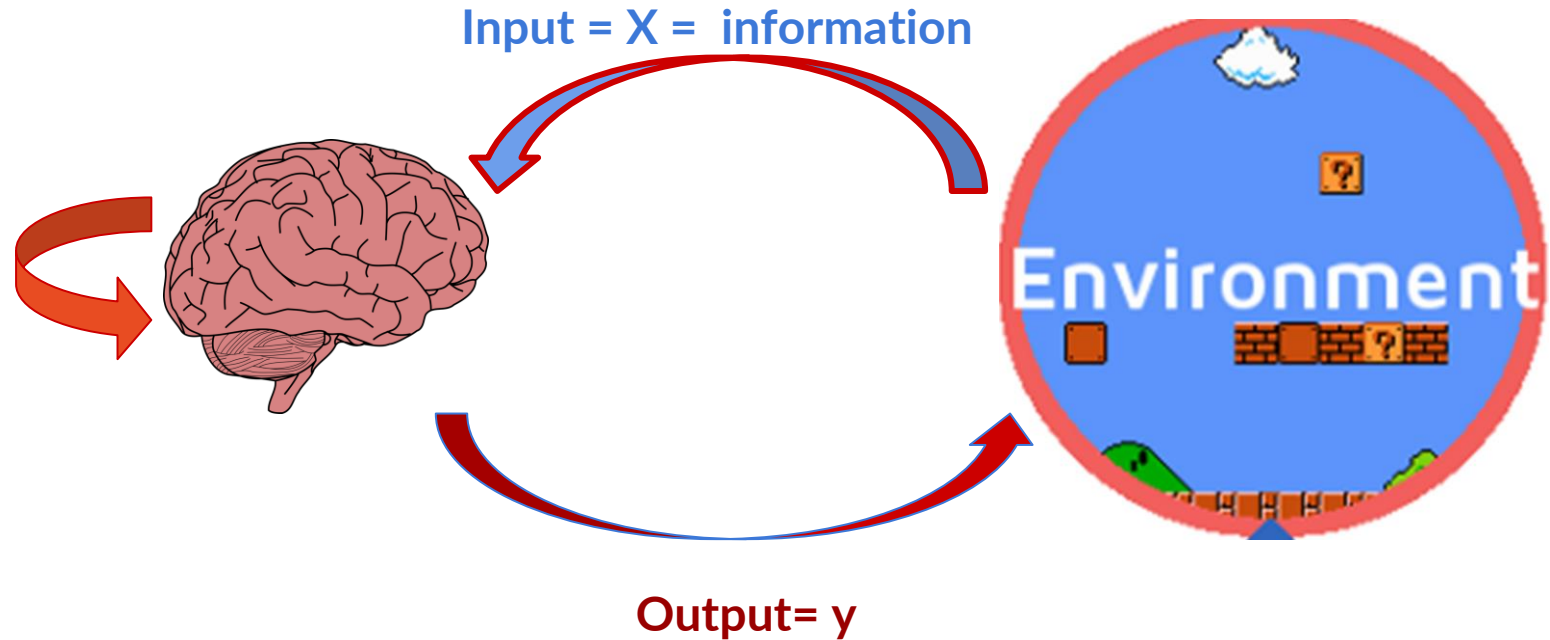
---

# Neurons and Artificial Neural Networks



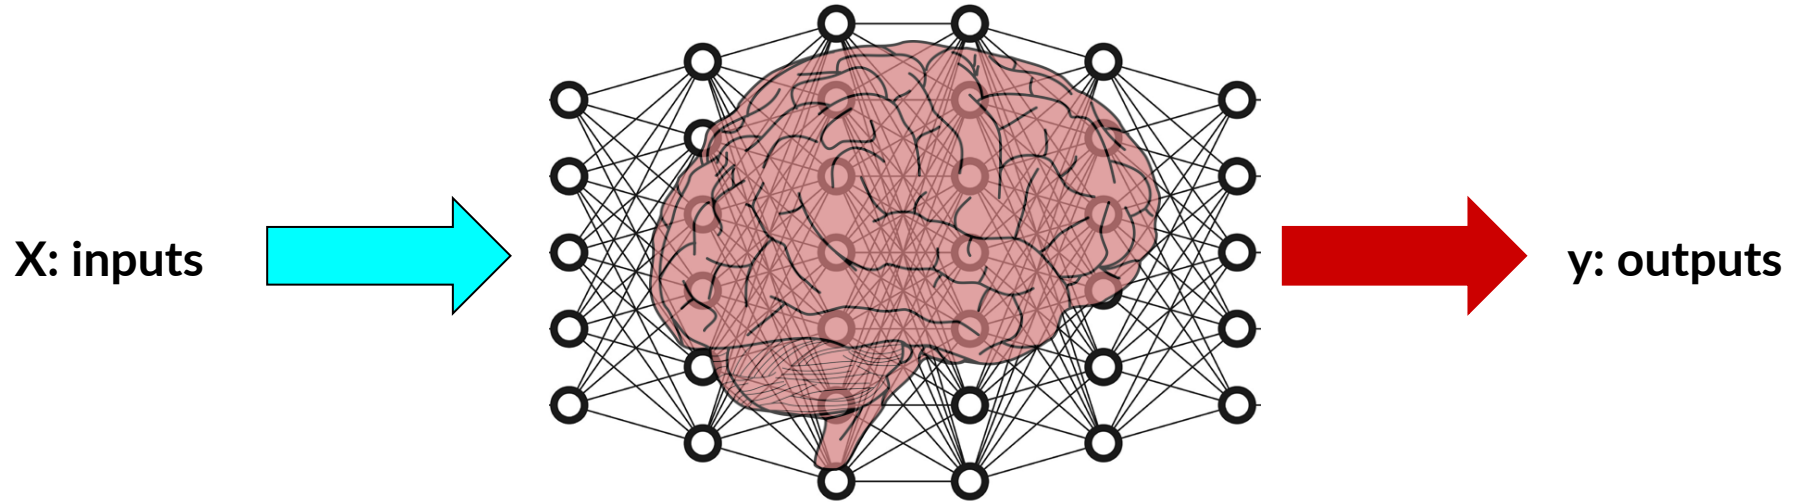
---

# Neurons and Artificial Neural Networks



---


# Neurons and Artificial Neural Networks



---

# Artificial Neural Networks can be classifiers

Dataset



60,000 small square 28x28 pixel grayscale **images of handwritten single digits between 0 and 9**

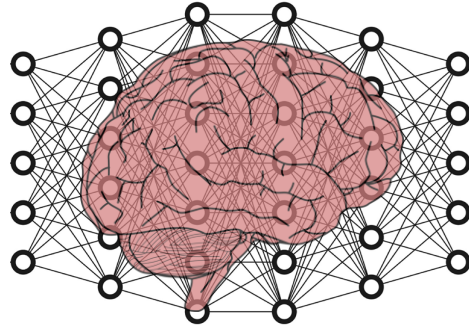
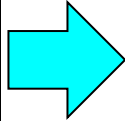
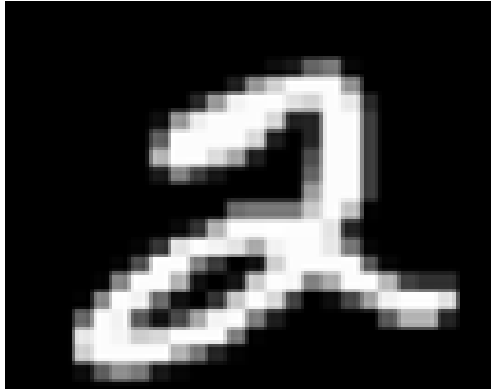


---

# Artificial Neural Networks can be classifiers

Image recognition

X:


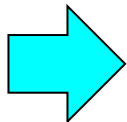
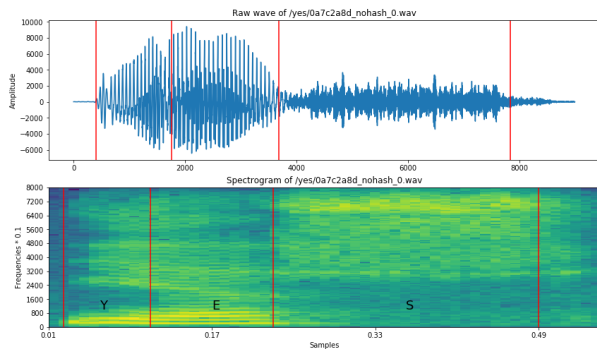


y: 2

# Artificial Neural Networks can be classifiers

## Speech recognition

X:



y: "yes"

Spectrogram of the spoken word "yes"

# Artificial Neural Networks

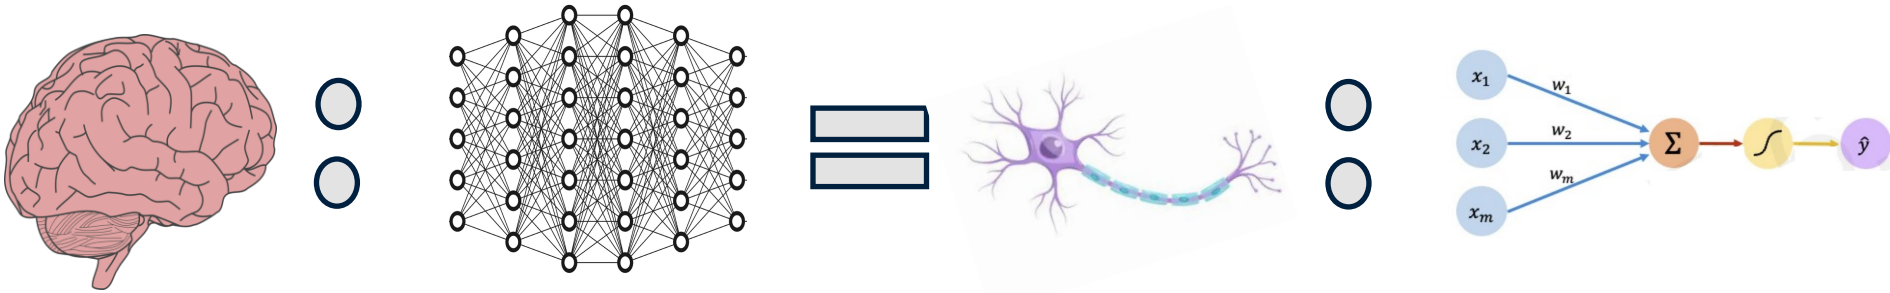
---

Can do countless things, including :

- **Natural Language Processing (NLP):**  
**Input:** Text data (e.g., a sentence or document).  
**Output:** Processed text (e.g., sentiment score, translated text, or named entities).
  - **Medical Diagnosis:**  
**Input:** Medical images (e.g., X-rays, MRIs) or patient data (e.g., age, symptoms).  
**Output:** Diagnosis or probability of a condition (e.g., presence of a tumor).
  - **Financial Services:**  
**Input:** Historical financial data (e.g., stock prices, transaction records).  
**Output:** Predictions (e.g., future stock prices, fraud detection alerts).
  - **Autonomous Vehicles:**  
**Input:** Sensor data (e.g., camera images, LIDAR scans).  
**Output:** Driving actions (e.g., steering angle, acceleration, braking).
  - **Recommendation Systems:**  
**Input:** User behaviour data (e.g., past purchases, viewing history).  
**Output:** Recommendations (e.g., movies, products)
-

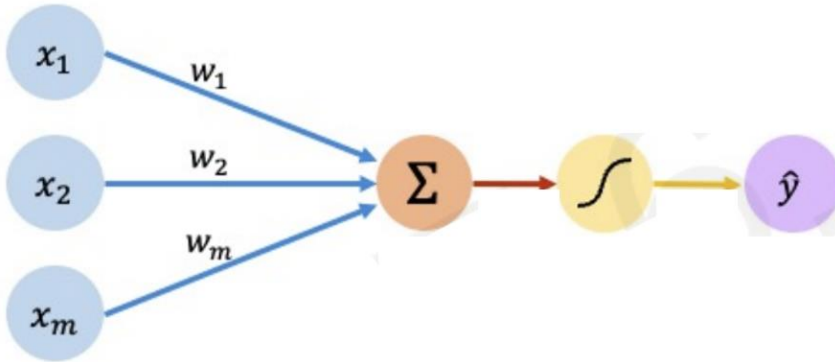
---

# Perceptron: the singular neuron



# Perceptron: the singular neuron

The building block of Artificial Neural Networks



Inputs    Weights    Sum    Non-Linearity    Output

A diagram illustrating the mathematical formula for a perceptron's output. The formula is  $\hat{y} = g \left( \sum_{i=1}^m x_i w_i \right)$ . A purple arrow points from the label "Output" to  $\hat{y}$ . A red arrow points from the label "Linear combination of inputs" to the summation term  $\sum_{i=1}^m x_i w_i$ . A yellow arrow points from the label "Non-linear activation function" to the function  $g$ .

## Does this remind you of anything?

# Logistic regression

There are many important research topics for which the dependent variable is "limited."

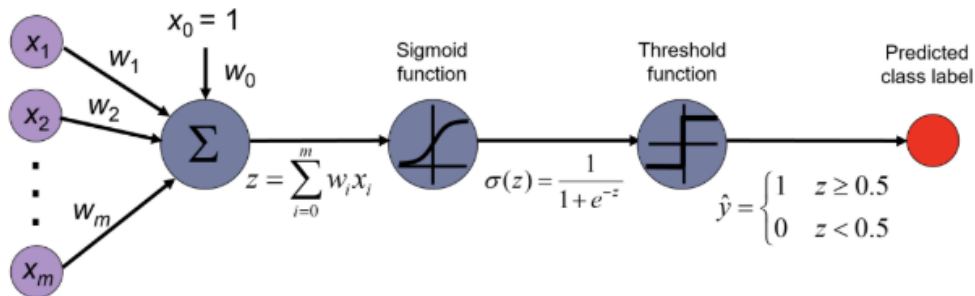
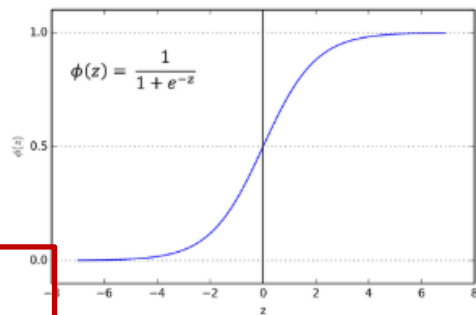
*For example:* voting, mortality, or number of accidents is not continuous or distributed normally.

$$Z = X \cdot W = w_0 + w_1 \cdot x_1 + w_2 \cdot x_2 + \dots + w_n \cdot x_n$$


Uses the sigmoid function to enforce the classification

You can also have a bias term

Number of weights = number of features + one bias



# Perceptron: the singular neuron



Output

Linear combination of inputs

$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$


Non-linear activation function

Bias

Inputs      Weights      Sum      Non-Linearity      Output

The bias term is a scalar that allows us to shift left or right along our activation function

# Perceptron: the singular neuron



Output

Linear combination of inputs

$$\hat{y} = g \left( w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function

Bias

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$

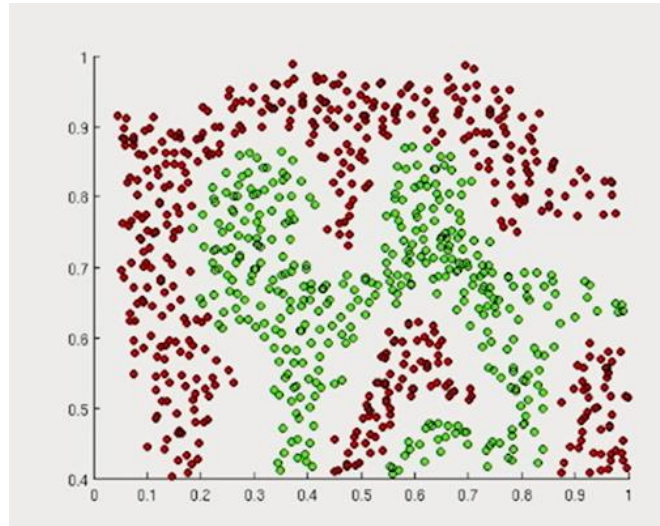
$$\text{where: } \mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix} \text{ and } \mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$$



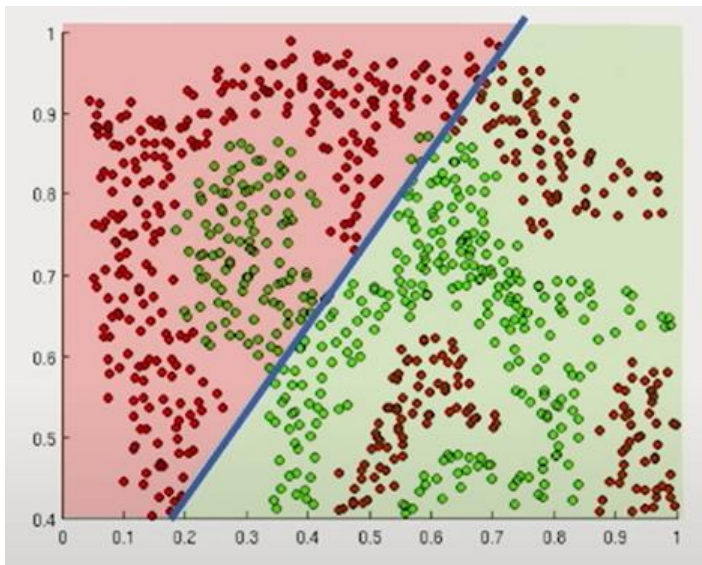
---

# Why do we need Activation Functions?

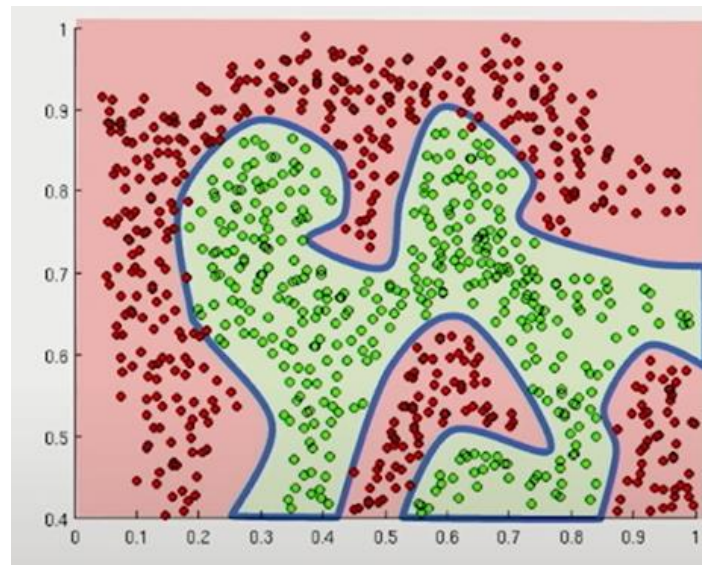
The purpose of activation functions is to introduce **non-linearities** into the network



# Why do we need non-linearity?



**Linear activation functions** produce **linear decisions** no matter the network size

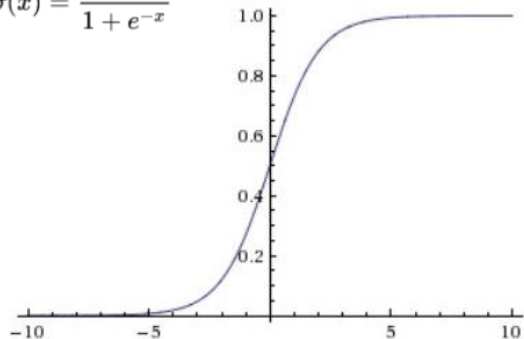


**Non-linearities** allow us to **approximate** arbitrarily complex functions

---

# What types of activation functions are there and how do we pick one?

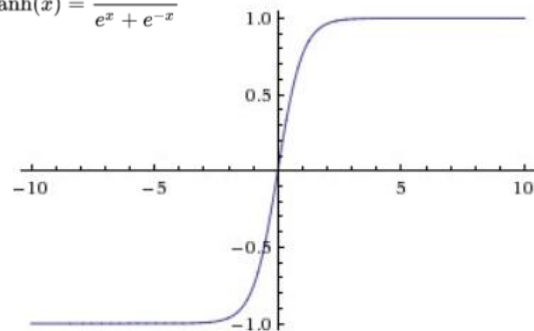
$$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$$



**Sigmoid** non-linearity squashes the numbers to range between [0,1].

Used for models where we have to predict the probability as an output

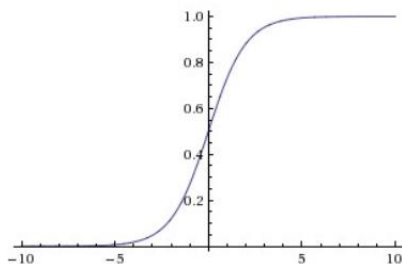
$$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



**tanh** non-linearity squashes the numbers to range between [-1,1]

Tanh tends to be preferred to sigmoid

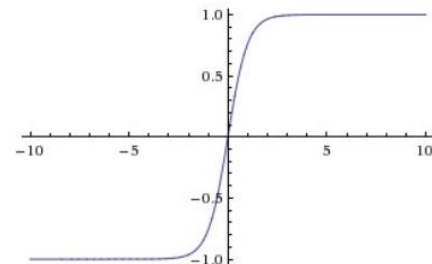
# What types of activation functions are there and how do we pick one?



```
from sklearn.neural_network import MLPClassifier

# Define the neural network with the sigmoid activation
function
clf = MLPClassifier(activation='logistic', hidden_layer_sizes=
(100,), max_iter=300)

# Train the neural network on the training data
clf.fit(X_train, y_train)
```



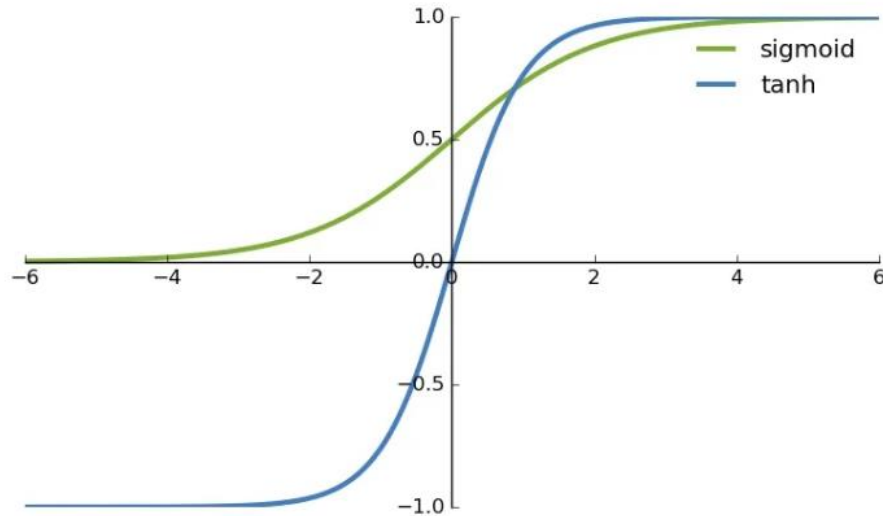
```
from sklearn.neural_network import MLPClassifier

# Define the neural network with the tanh activation function
clf = MLPClassifier(activation='tanh', hidden_layer_sizes=
(100,), max_iter=300)

# Train the neural network on the training data
clf.fit(X_train, y_train)
```

---


# What types of activation functions are there and how do we pick one?



---

# What types of activation functions are there and how do we pick one?

$$f(x) = \max(0, x)$$



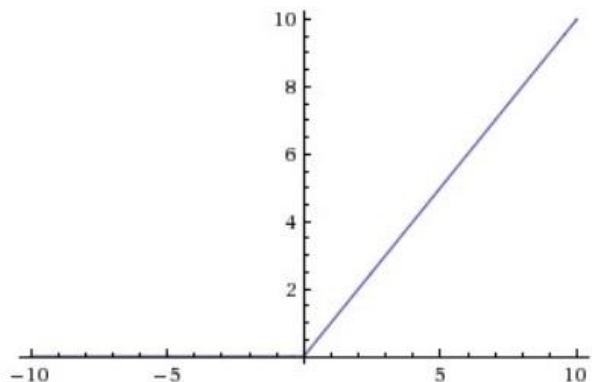
**ReLU** (Rectified Linear Unit) activation function is zero when  $x < 0$  and then linear with slope 1 when  $x > 0$

- **ReLU speeds up** the calculations up to 6x compared to **tanh**
- **ReLU units** can be fragile during training and can “die”. (i.e. neurons that never activate across the entire training dataset) if the learning rate is set too high. **With a proper setting of the learning rate this is less frequently an issue.**

---

# What types of activation functions are there and how do we pick one?

$$f(x) = \max(0, x)$$



**ReLU** (Rectified Linear Unit) activation function is zero when  $x < 0$  and then linear with slope 1 when  $x > 0$

```
from sklearn.neural_network import MLPClassifier


# Define the neural network with the ReLU activation function
clf = MLPClassifier(activation='relu', hidden_layer_sizes=(100,), max_iter=300)

# Train the neural network on the training data
clf.fit(X_train, y_train)
```

---

# What types of activation functions are there and how do we pick one?

$$f(x) = \begin{cases} 0.01x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

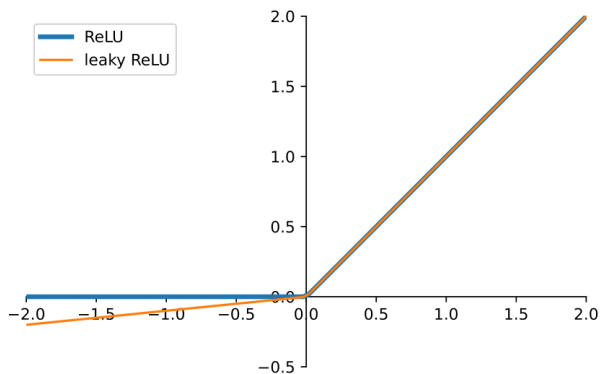


- **Leaky ReLUs** are one attempt to fix the “dying ReLU” problem.
- Instead of the function being zero when  $x < 0$ , a leaky ReLU will instead have a small positive slope (of 0.01, or so). That is, the function computes  $f(x) = 1(x < 0)(\alpha x) + 1(x \geq 0)(x)$  where  $\alpha$  is a small constant.

Other solutions to the “dying ReLU” problem include ELUs (Exponential Linear Units) and PReLUs (Parametric ReLU).



# What types of activation functions are there and how do we pick one?



- **Leaky ReLUs** are one attempt to fix the “dying ReLU” problem.

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LeakyReLU


# Define the neural network model
model = Sequential([
    # Add a dense layer with 100 neurons
    Dense(100, input_dim=input_dim),
    # Apply Leaky ReLU activation function with alpha=0.01
    LeakyReLU(alpha=0.01),
    # Add the output layer with softmax activation for multi-class classification
    Dense(output_dim, activation='softmax')
])

# Compile the model with Adam optimizer and categorical crossentropy loss
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

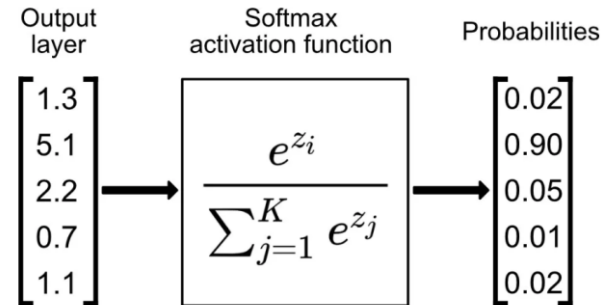
# Train the neural network on the training data
model.fit(X_train, y_train, epochs=300)
```

---






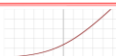


# What types of activation functions are there and how do we pick one?



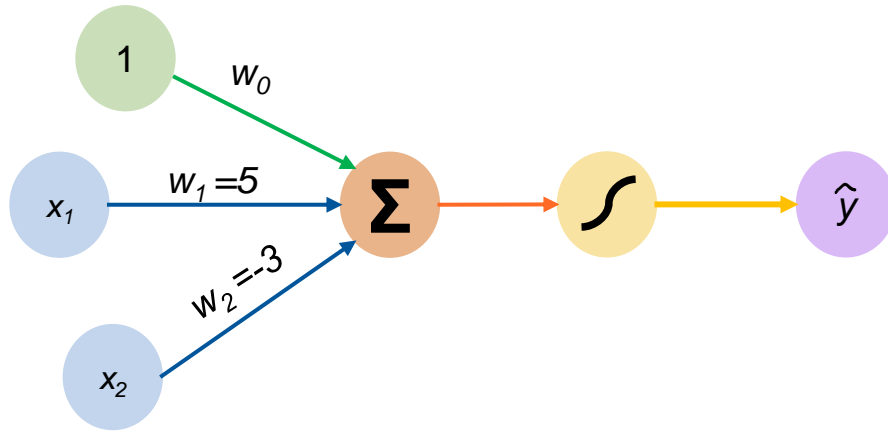
- **Softmax** is used when you need to classify data into more than two categories. For example, if you have a dataset with images of cats, dogs, and birds, Softmax can help determine the probability that a given image belongs to each of these classes.
- It is applied after the output.



# What types of activation functions are there and how do we pick one?

ACTIVATION FUNCTION	PLOT	EQUATION	DERIVATIVE	RANGE
Linear		$f(x) = x$	$f'(x) = 1$	$(-\infty, \infty)$
Binary Step		$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{if } x \neq 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$	$\{0, 1\}$
Sigmoid		$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$	$(0, 1)$
Hyperbolic Tangent(tanh)		$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	$f'(x) = 1 - f(x)^2$	$(-1, 1)$
Rectified Linear Unit(ReLU)		$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$	$[0, \infty)$
Softplus		$f(x) = \ln(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$	$(0, 1)$
Leaky ReLU		$f(x) = \begin{cases} 0.01x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0.01 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	$(-1, 1)$
Exponential Linear Unit(ELU)		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$	$f'(x) = \begin{cases} \alpha e^x & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ 1 & \text{if } x = 0 \text{ and } \alpha = 1 \end{cases}$	$[0, \infty)$

# Perceptron: an example



Inputs

Weights

Sum

Non-Linearity

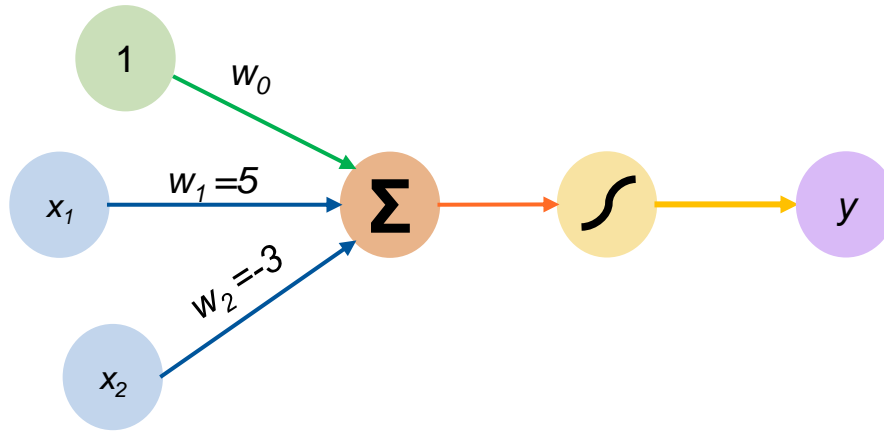
Output

We have  $w_0 = 1$  and  $\mathbf{W} = \begin{bmatrix} 5 \\ -3 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{W}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 5 \\ -3 \end{bmatrix}\right) \\ \hat{y} &= g(1 + \underbrace{5x_1 - 3x_2})\end{aligned}$$

This is just a line in 2D!

# Perceptron: an example



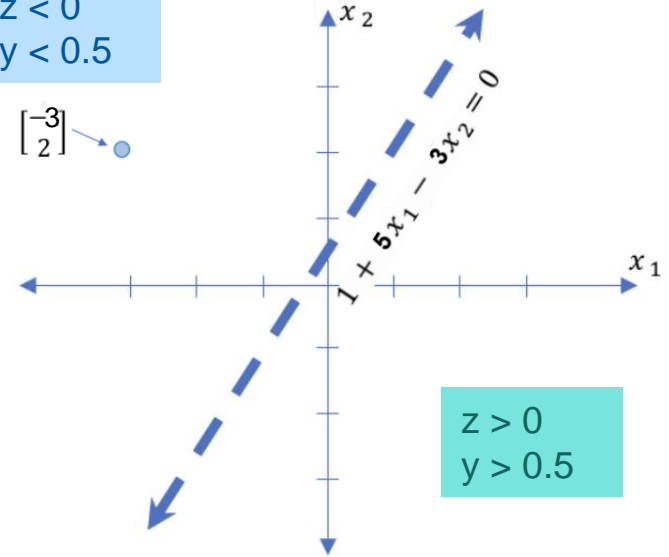
If we have an input  $\mathbf{X} = \begin{bmatrix} -3 \\ 2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(1 + (5 \cdot -3) - (3 \cdot 2)) \\ &= g(-20) = 2.06 \text{ E-9}\end{aligned}$$

$$\hat{y} = g(\underbrace{1 + 5x_1 - 3x_2}_{(z)})$$

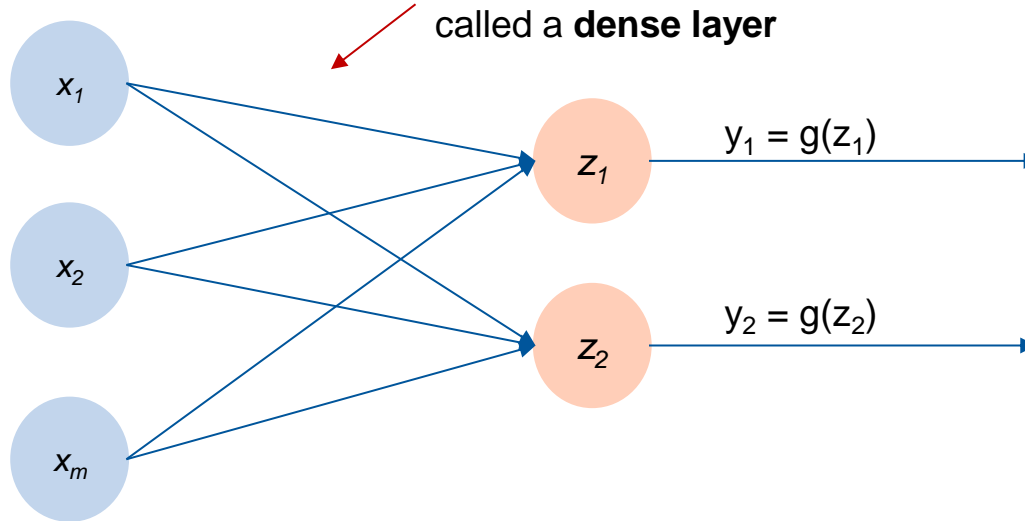
$$\begin{aligned}z &< 0 \\ y &< 0.5\end{aligned}$$

$$\begin{bmatrix} -3 \\ 2 \end{bmatrix}$$



# Multi-output perceptron

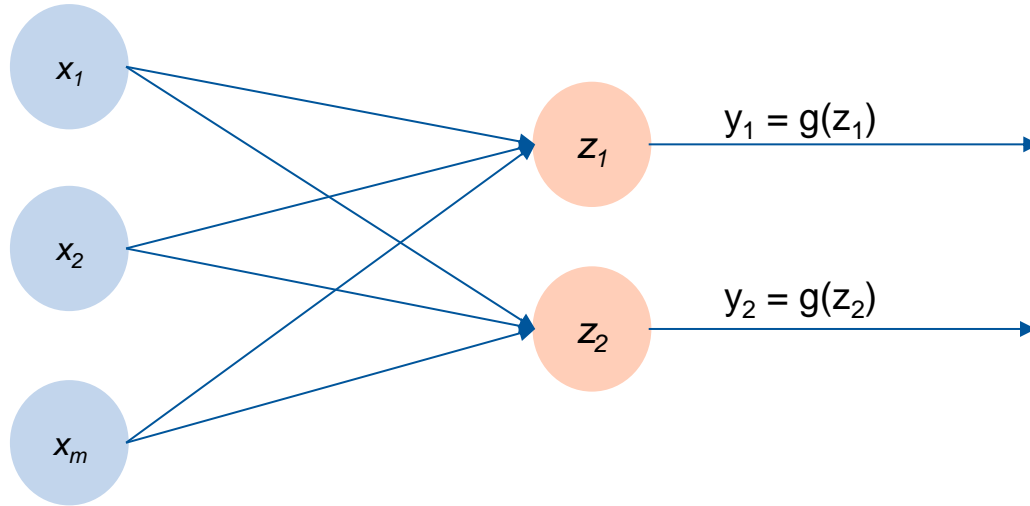
Because of the dense connections between inputs and outputs, this is called a **dense layer**



This is just an example to understand the intuition, but multi-output perceptrons are useful when you need to predict multiple target variables simultaneously. For example, predicting both temperature and humidity based on weather data.

$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

# Multi-output perceptron



```
from sklearn.neural_network import MLPRegressor
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split

# Generate synthetic data with multiple outputs
X, y = make_regression(n_samples=1000, n_features=20,
                      n_targets=2, noise=0.1)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

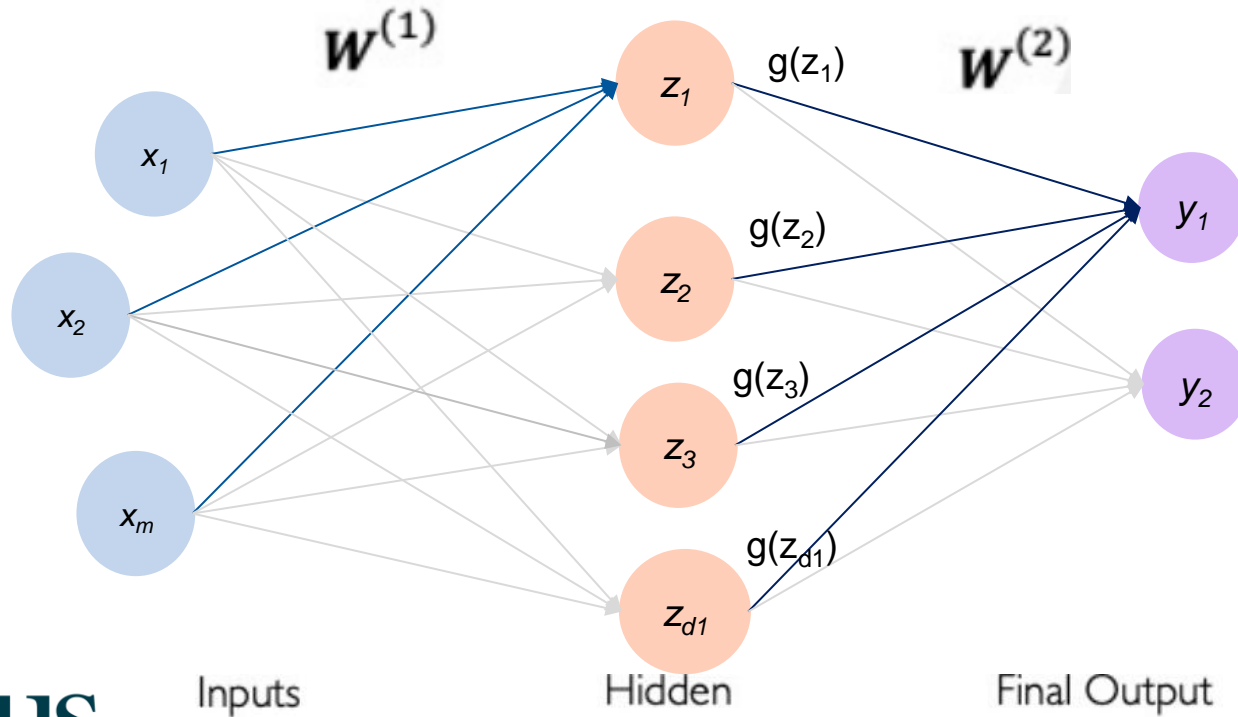
# Define the multi-output perceptron
mlp = MLPRegressor(hidden_layer_sizes=(100,),
                  activation='relu', max_iter=300)

# Train the neural network on the training data
mlp.fit(X_train, y_train)

# Make predictions on the test data
predictions = mlp.predict(X_test)
```

$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

# Single Layer Neural Network



In layered neural networks, the outputs from one layer become the inputs of the next layer



---

# Single Layer Neural Network

For classification:

```
from sklearn.neural_network import MLPClassifier
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split

# Generate synthetic classification data
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2,
random_state=42)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define the single-layer neural network
clf = MLPClassifier(hidden_layer_sizes=(100,), activation='relu',
max_iter=300)

# Train the neural network on the training data
clf.fit(X_train, y_train)

# Make predictions on the test data
predictions = clf.predict(X_test)
```

For regression:

```
from sklearn.neural_network import MLPRegressor
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split

# Generate synthetic regression data
X, y = make_regression(n_samples=1000, n_features=20, noise=0.1,
random_state=42)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define the single-layer neural network
reg = MLPRegressor(hidden_layer_sizes=(100,), activation='relu',
max_iter=300)

# Train the neural network on the training data
reg.fit(X_train, y_train)

# Make predictions on the test data
predictions = reg.predict(X_test)
```

---

# Single Layer Neural Network

For classification:

```
from sklearn.neural_network import MLPClassifier
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split

# Generate synthetic classification data
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2,
random_state=42)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define the single-layer neural network
clf = MLPClassifier(hidden_layer_sizes=(100,), activation='relu',
max_iter=300)

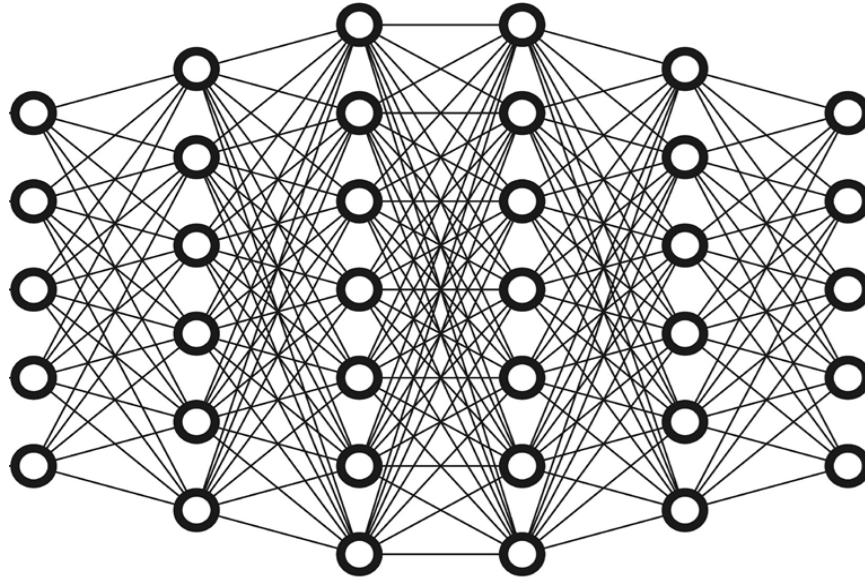
# Train the neural network on the training data
clf.fit(X_train, y_train)

# Make predictions on the test data
predictions = clf.predict(X_test)
```

- **n\_samples**= The number of samples (data points) in your dataset. 1000 creates 1000 data points.
- **n\_features**=The number of features (input variables) in your dataset. 20 means each data point has 20 input variables.
- **n\_classes**= It determines how many different categories the target variable (y) can take. 2 means the target variable (y) will have three distinct class labels (e.g., 0, 1).
- **random\_state**=A seed value for random number generation to ensure reproducibility. 42 ensures the data generation process is reproducible.

---

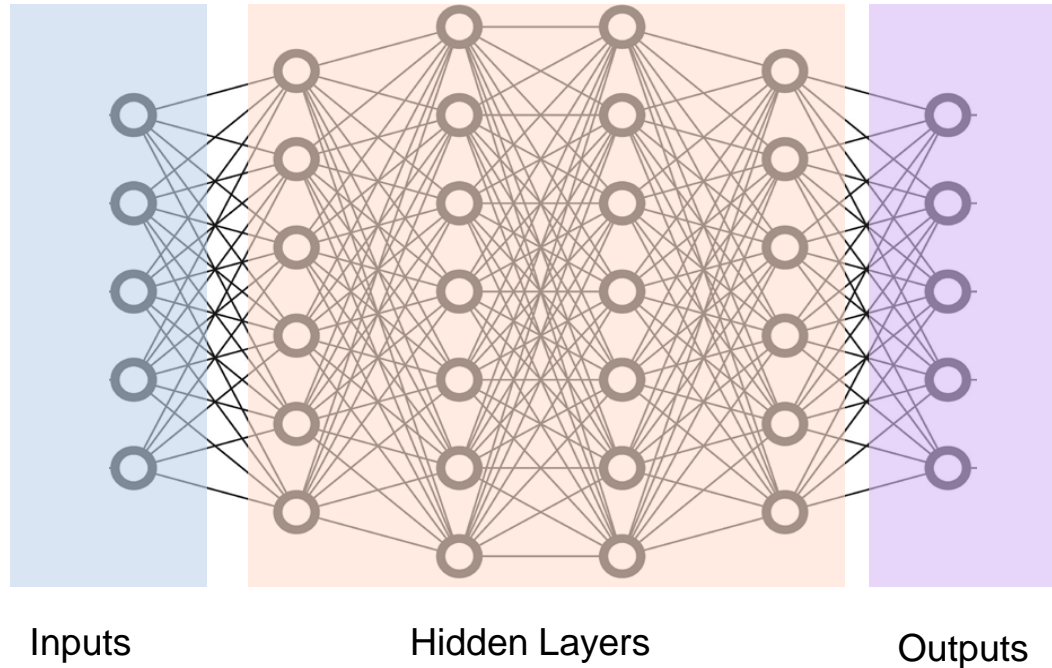
# Multi Layer Neural Network



In layered neural networks, the outputs from one layer become the inputs of the next layer

---

# Multi Layer Neural Network



In layered neural networks, the outputs from one layer become the inputs of the next layer

---

# Multi Layer Neural Network

For classification:

```
from sklearn.neural_network import MLPClassifier
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split

# Generate synthetic classification data
# n_samples: number of samples
# n_features: number of features
# n_classes: number of classes
# random_state: seed for reproducibility
X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)

# Split the data into training and testing sets
# test_size: proportion of the dataset to include in the test split
# random_state: seed for reproducibility
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define the multi-layer neural network with two hidden layers
# hidden_layer_sizes: tuple specifying the number of neurons in each hidden layer
# activation: activation function for the hidden layers ('relu' in this case)
# max_iter: maximum number of iterations for training
clf = MLPClassifier(hidden_layer_sizes=(100, 50), activation='relu', max_iter=300)

# Train the neural network on the training data
# X_train: training input data
# y_train: training target data
clf.fit(X_train, y_train)

# Make predictions on the test data
# X_test: testing input data
predictions = clf.predict(X_test)
```

- **hidden\_layer\_sizes=** in this case it specifies two hidden layers, the first with 100 neurons and the second with 50 neurons.

---

---

# Training a neural network


---

---

**The code does it for us. But what does it do exactly?**

---

# Think about your data and think about this model



## Is there any information we don't have?

---



---

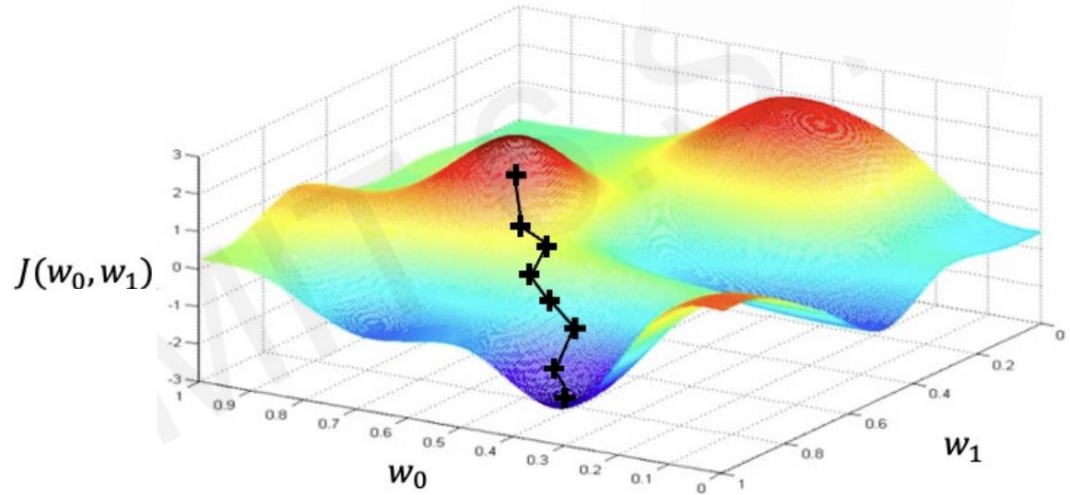
# Loss optimisation

- Quantify the Loss (how wrong the outputs  $y$  are compared to our test data)
- Update the weights (through gradient descent and backpropagation) to minimise this loss

---

# Loss optimisation

- Quantify the Loss (how wrong the outputs  $y$  are compared to our test data)
- Update the weights (through gradient descent and backpropagation) to minimise this loss



---

---

scikit-learn handles all the underlying details of backpropagation and gradient descent, allowing you to focus on defining and training your model

**The code handles this for us**

---

---

# Types of ANNs

---

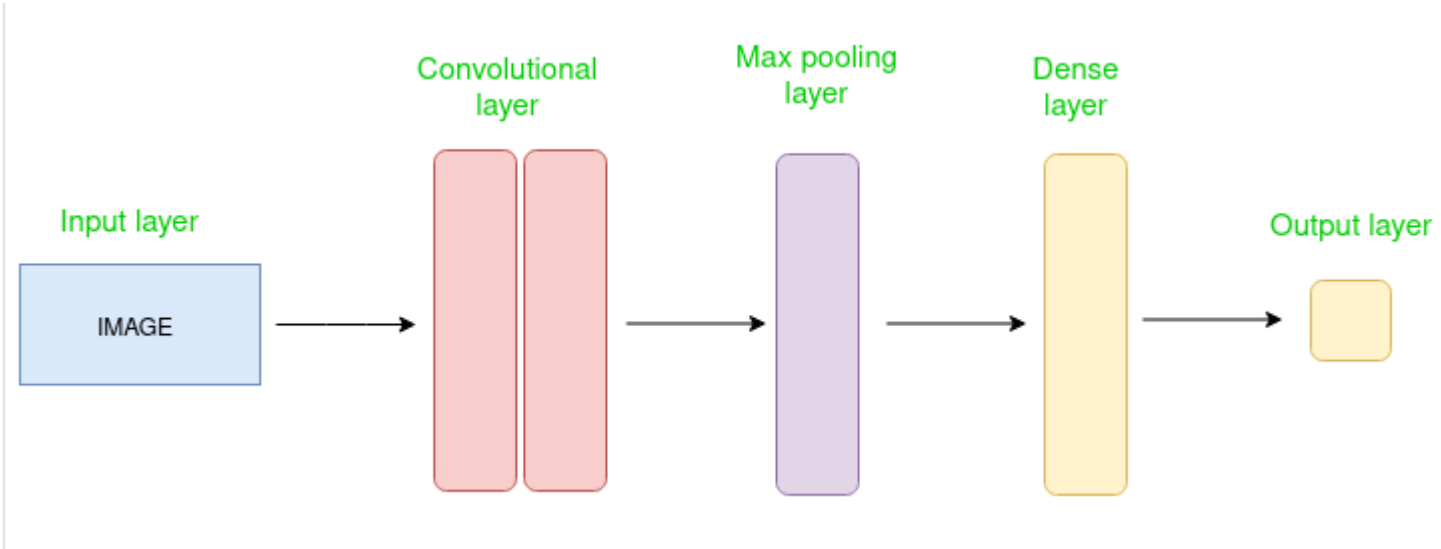
# Convolutional Neural Networks (CNNs)

CNNs are designed for visual data processing.



---

# Convolutional neural networks



---

# Convolutional neural networks

Input layer

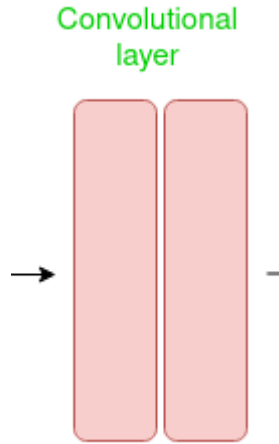


Let's take an example by running an image of dimension  $32 \times 32 \times 3$ .

- **Input Layers:** It's the layer in which we give input to our model. In CNN, Generally, the input will be an image or a sequence of images. This layer holds the raw input of the image with width 32, height 32, and depth 3 (= the colour array).

---

# Convolutional neural networks

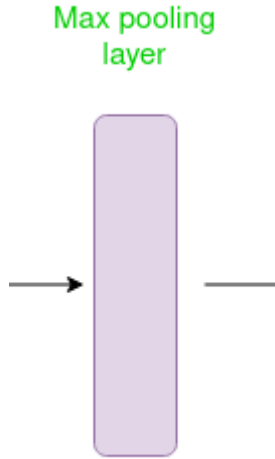


- **Filters/Kernels:** Small matrices that slide over the input image. Each filter is designed to detect specific features like edges, textures, or patterns. Imagine looking at an image through a small window that slides over the picture. Each window (filter) looks for specific things like edges or patterns.
- **Convolution Operation:** The filter is applied to the input image, producing a feature map. This operation involves multiplying the filter values with the corresponding pixel values and summing them up. The result is a new matrix that highlights the presence of the feature detected by the filter. As the window slides over the image, it creates a new picture (feature map) that highlights where it found those edges or patterns.
- **ReLU:** After the window slides over the image, ReLU changes all negative values to zero. This helps the network focus on important features and ignore less useful information.



---

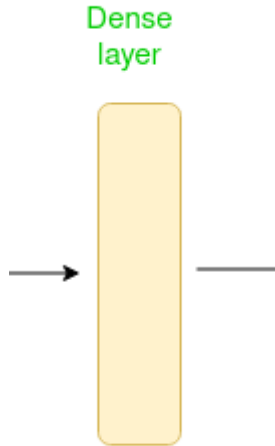
# Convolutional neural networks



- **Max Pooling:** This layer reduces the dimensionality of the feature map while retaining important features. It does this by selecting the maximum value from a small region (e.g., 2x2) of the feature map. This helps in reducing computational load and preventing overfitting.
- Think of zooming out on a picture to make it smaller while keeping the important parts. Max pooling picks the biggest value from a small area, reducing the size of the feature map but keeping the key features.

---

# Convolutional neural networks



- **Flattening:** Converts the 2D feature maps into a 1D vector. This step prepares the data for the fully connected layers. **Imagine taking all the important features from the picture and laying them out in a single line.**
- **Dense Layers:** These layers connect every neuron from one layer to every neuron in the next layer. They combine the features learned by the convolutional and pooling layers to make final predictions. **This is where the network decides which class the input image belongs to.**

---

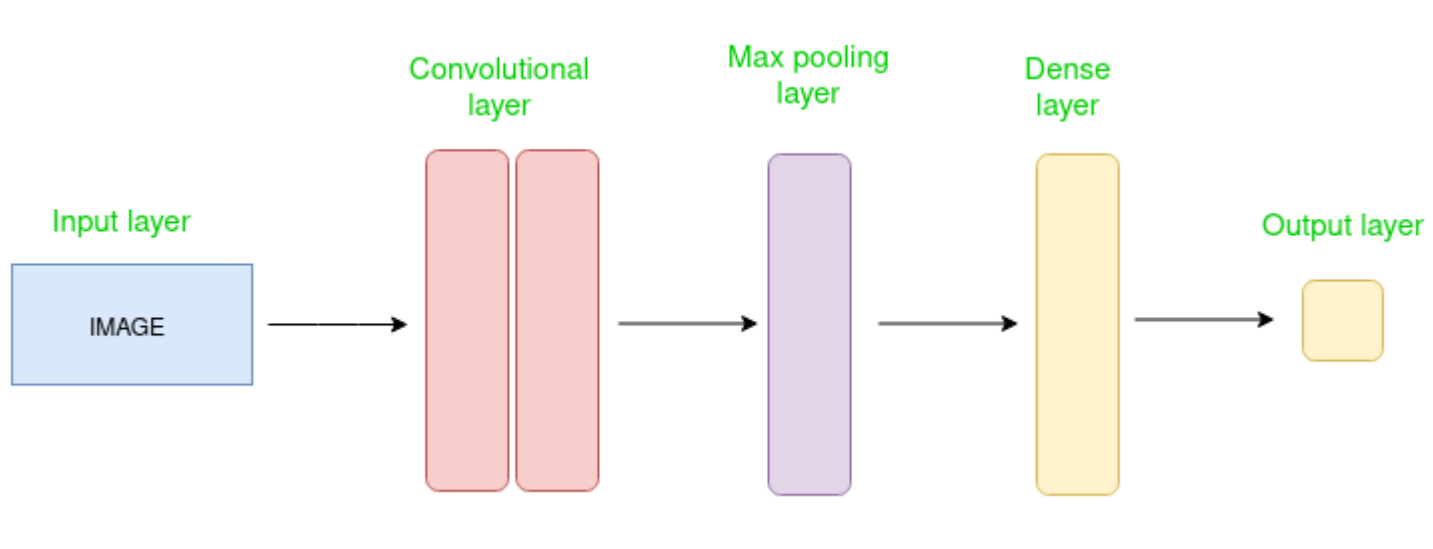
# Convolutional neural networks



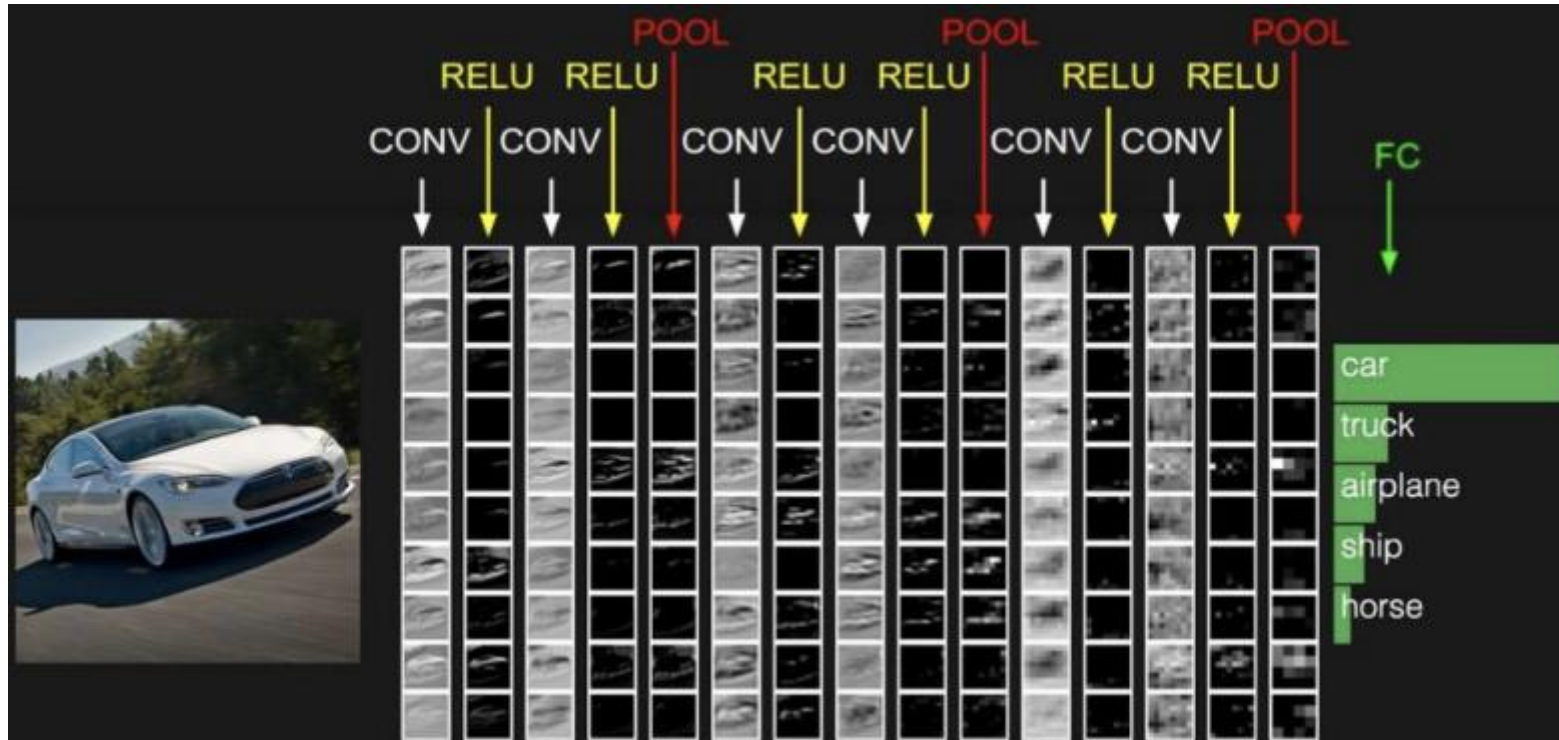
- **Output Layer:** The output from the fully connected layers is then fed into a logistic function for classification tasks like sigmoid or softmax which converts the output of each class into the probability score of each class.

---

# Convolutional neural networks



# Convolutional neural networks



# Convolutional neural networks

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

# Define the CNN model
model = Sequential()

# Add a convolutional layer with 32 filters, a 3x3 kernel, and ReLU activation
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(64, 64, 3)))

# Add a max pooling layer with a 2x2 pool size
model.add(MaxPooling2D(pool_size=(2, 2)))

# Add another convolutional layer with 64 filters
model.add(Conv2D(64, (3, 3), activation='relu'))

# Add another max pooling layer
model.add(MaxPooling2D(pool_size=(2, 2)))

# Flatten the feature maps to a 1D vector
model.add(Flatten())

# Add a fully connected layer with 128 neurons and ReLU activation
model.add(Dense(128, activation='relu'))

# Add the output layer with softmax activation for multi-class classification
model.add(Dense(10, activation='softmax'))

# Compile the model with Adam optimizer and categorical crossentropy loss
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Print the model summary
model.summary()
```

## 1.Convolutional Layers:

1. **Number of Filters:** The choice of 32 and 64 filters is typical for initial layers in a CNN. These numbers can be adjusted based on the complexity of the task and the size of the dataset.
2. **Kernel Size:** A 3x3 kernel is a standard choice that balances computational efficiency and the ability to capture spatial features.

## 2.Pooling Layers:

1. **Pool Size:** A 2x2 pool size is commonly used to reduce the spatial dimensions while retaining important features.

## 3.Fully Connected Layers:

1. **Number of Neurons:** The choice of 128 neurons in the dense layer is a common practice, but it can be tuned based on the complexity of the task.

## 4.Activation Functions:

1. **ReLU:** The ReLU activation function is widely used due to its ability to mitigate the vanishing gradient problem and speed up training.
2. **Softmax:** Softmax is used in the output layer for multi-class classification tasks.

## 5.Optimizer and Loss Function:

1. **Adam Optimizer:** Adam is a popular choice for its adaptive learning rate and efficiency.
2. **Categorical Crossentropy:** This loss function is suitable for multi-class classification problems.

---

# Recursive Neural Networks (RNNs)

Language  
Modeling  
and Text  
Generation

Speech  
Recognition

Sentiment  
Analysis

Music  
Composition

Machine  
Translation

Time Series  
Prediction


Video  
Analysis

Healthcare

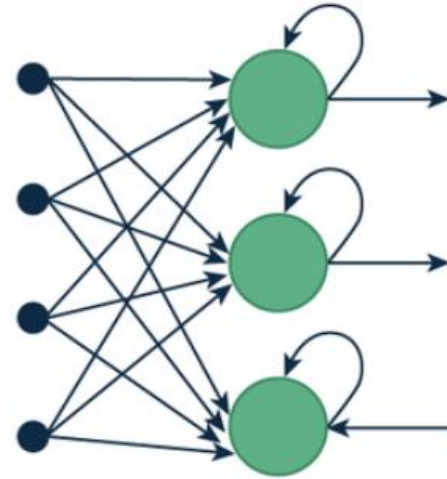
---

# Recursive Neural Networks (RNNs)

RNNs are unique because they have **loops in their architecture**, allowing information to persist. This makes them well-suited for **tasks where context or previous information is important**, such as language modeling or predicting stock prices.



Feed-Forward Neural Network



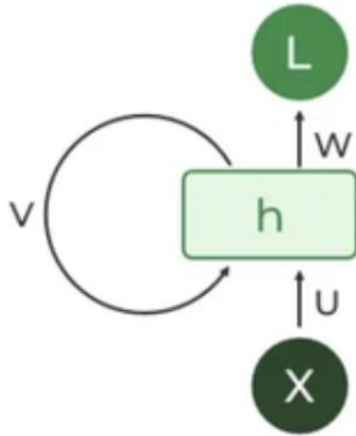
Recurrent Neural Network



---

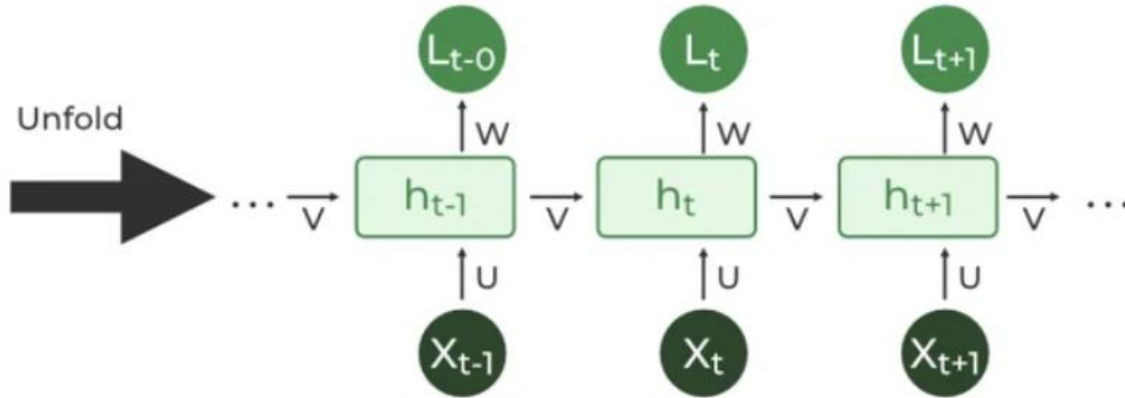
# Recursive Neural Networks (RNNs)

The fundamental processing unit in RNN is a **Recurrent Unit**.



- Recurrent units hold a **hidden state that maintains information about previous inputs in a sequence**. Recurrent units can “remember” information from prior steps by feeding back **their hidden state**

# Recursive Neural Networks (RNNs)

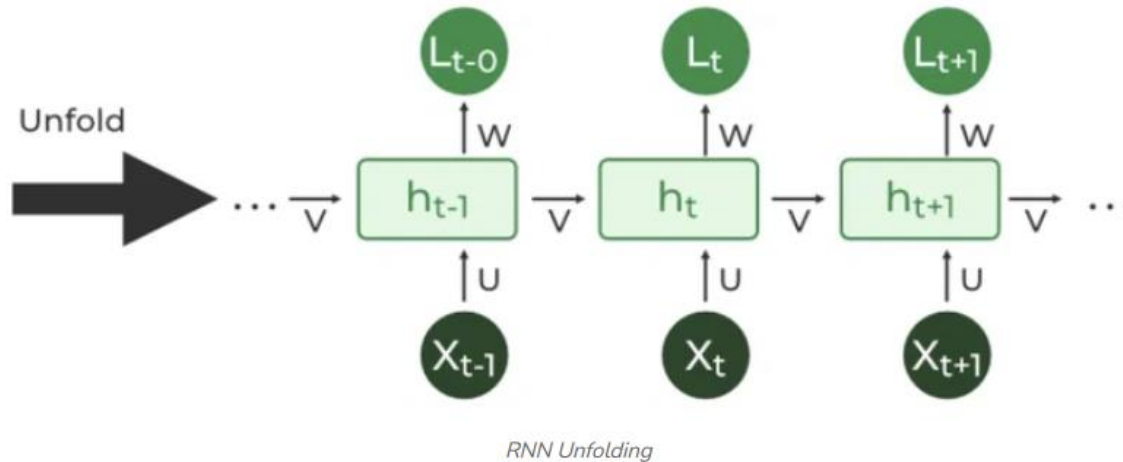


*RNN Unfolding*

- **Sequence Processing:** RNNs process input sequences one element at a time, updating the hidden state at each step.

---

# Recursive Neural Networks (RNNs)

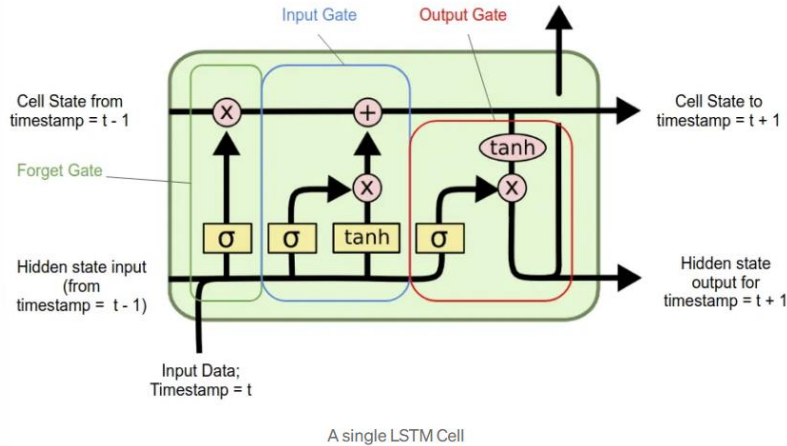


Since RNNs process sequential data [Backpropagation Through Time \(BPTT\)](#) is used to update the network's parameters.

But this is also handled by code, so we don't need to worry about this!

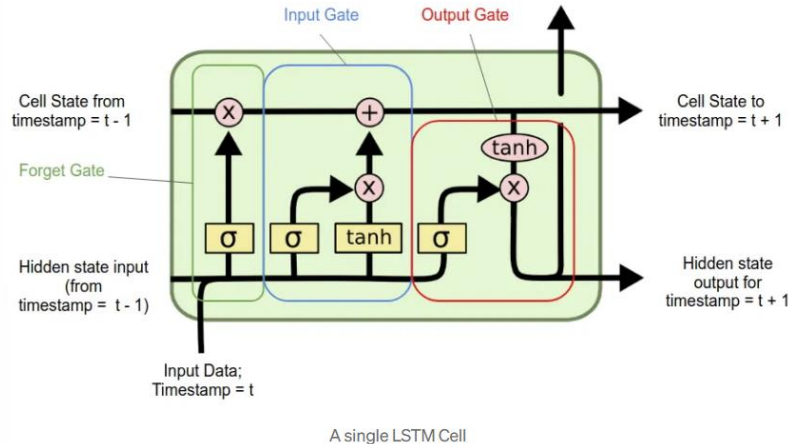
# Long Short-Term Memory (LSTM) networks

LSTMs are a type of Recurrent Neural Network (RNN) designed to address the limitations of traditional RNNs, particularly the issue of long-term dependencies



- **Cell State:** LSTMs have a cell state that runs through the entire sequence, allowing information to be carried over long distances.
- **Gates:** LSTMs use gates to control the flow of information. These gates are:
  - **Forget Gate:** Decides what information to discard from the cell state.
  - **Input Gate:** Decides what new information to add to the cell state.
  - **Output Gate:** Decides what information to output from the cell state.

# Long Short-Term Memory (LSTM) networks



## 1. Forget Gate:

- The forget gate determines which parts of the cell state to forget. It uses a sigmoid function to output values between 0 and 1, where 0 means "forget completely" and 1 means "keep completely."

## 2. Input Gate:

- The input gate decides which new information to add to the cell state. It uses a sigmoid function to determine which values to update and a tanh function to create new candidate values.

## 3. Cell State Update:

- The cell state is updated by combining the forget gate's output and the input gate's new candidate values.

## 4. Output Gate:

- The output gate determines the final output of the LSTM cell. It uses a sigmoid function to decide which parts of the cell state to output and a tanh function to scale the output.

---

# RNNs implementations

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, Dense

# Define the RNN model
model = Sequential()

# Add a SimpleRNN layer with 50 units
# 50 units is a common choice for capturing patterns without overfitting
model.add(SimpleRNN(50, input_shape=(None, 1)))

# Add a fully connected layer with 1 neuron and sigmoid activation
# Sigmoid activation is used for binary classification
model.add(Dense(1, activation='sigmoid'))

# Compile the model with Adam optimizer and binary crossentropy loss
# Adam optimizer is efficient and adaptive, binary crossentropy is suitable for binary classification
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Print the model summary
model.summary()
```

**RNNs:** Suitable for tasks with short-term dependencies, such as simple sequence prediction.

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense

# Define the LSTM model
model = Sequential()

# Add an LSTM layer with 50 units
# LSTM units are optimal for capturing long-term dependencies
model.add(LSTM(50, input_shape=(None, 1)))

# Add a fully connected layer with 1 neuron and sigmoid activation
# Sigmoid activation is used for binary classification
model.add(Dense(1, activation='sigmoid'))

# Compile the model with Adam optimizer and binary crossentropy loss
# Adam optimizer is efficient and adaptive, binary crossentropy is suitable for binary classification
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Print the model summary
model.summary()
```

**LSTMs:** Suitable for tasks with long-term dependencies, such as language modeling, speech recognition, and time series forecasting.

---

# Recap

- What are Neural Networks
- Artificial Neurons (perceptron, MLPs, Single and multi layered NNs)
- Training Neural Networks
- Some types of Neural Networks (CNNs, RNNs, LSTMs)

---

---

# LAB



---

# Neural Nets Lab

---


- Go to github for lab8.ipynb

---

---

# Supplementary Material

# Biological vs Artificial neuron



A cartoon drawing of a biological neuron (left) and its mathematical model (right).

# Quantifying loss

$$\mathcal{L}(\underbrace{f(x^{(i)}; \mathbf{w})}_{\text{Predicted}}, \underbrace{y^{(i)}}_{\text{Actual}})$$

- **Empirical loss:** measures the total loss over the entire dataset

$f(x)$		$y$
0.1	✗	1
0.8	✗	0
0.6	✓	1
⋮		⋮

$$J(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(\underbrace{f(x^{(i)}; \mathbf{w})}_{\text{Predicted}}, \underbrace{y^{(i)}}_{\text{Actual}})$$

- **Binary Cross Entropy Loss:** can be used with models that output a probability between 0 and 1

$f(x)$		$y$
0.1	✗	1
0.8	✗	0
0.6	✓	1
⋮		⋮

$$J(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \underbrace{y^{(i)}}_{\text{Actual}} \log(\underbrace{f(x^{(i)}; \mathbf{w})}_{\text{Predicted}}) + (1 - \underbrace{y^{(i)}}_{\text{Actual}}) \log(1 - \underbrace{f(x^{(i)}; \mathbf{w})}_{\text{Predicted}})$$

- **Mean Squared Error Loss:** can be used with regression models that output continuous numbers

$f(x)$		$y$
30	✗	90
80	✗	20
85	✓	95
⋮		⋮

$$J(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n (\underbrace{y^{(i)}}_{\text{Actual}} - \underbrace{f(x^{(i)}; \mathbf{w})}_{\text{Predicted}})^2$$

---

# Loss optimisation

We want to find the network weights that **achieve the lowest loss**

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$

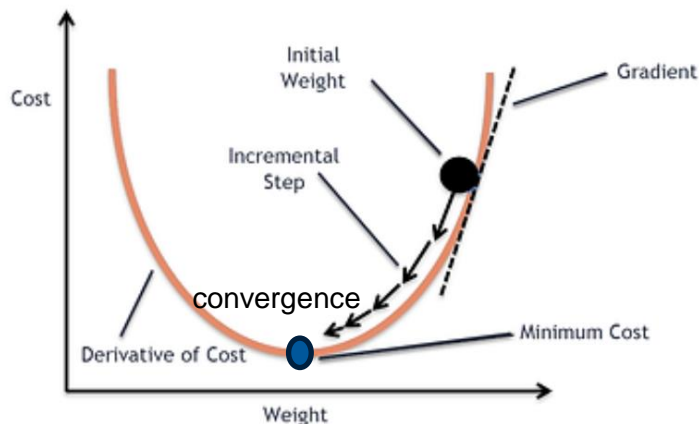
$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$

Remember:  
 $\mathbf{W} = \{\mathbf{w}^{(0)}, \mathbf{w}^{(1)}, \dots\}$

---

# Gradient descent

"A gradient measures how much the output of a function changes if you change the inputs a little bit." — Lex Fridman (MIT)


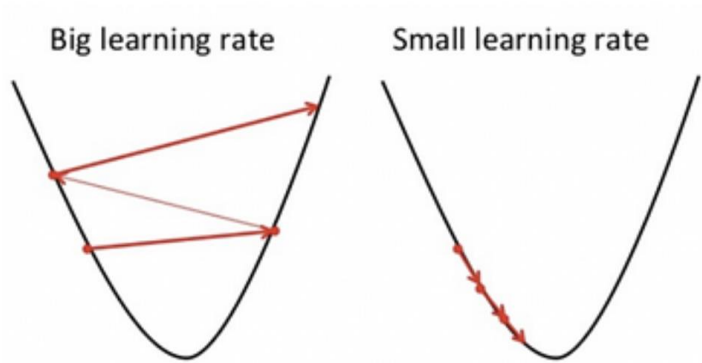


## Algorithm

1. Initialize weights randomly  $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3.   Compute gradient,  $\frac{\partial J(W)}{\partial W}$
4.   Update weights,  $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(W)}{\partial W}$
5. Return weights

---

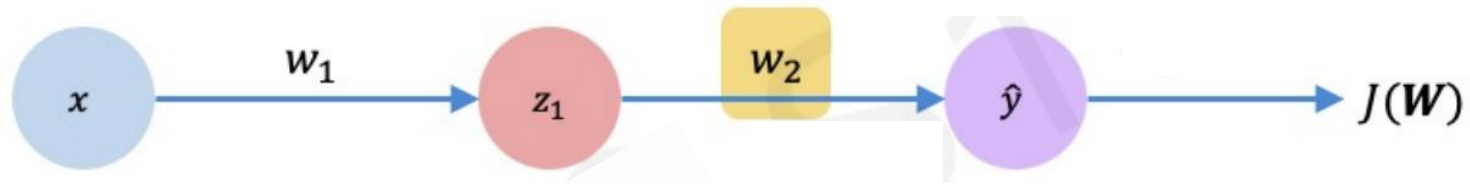
# Gradient Descent and learning rate



A big learning rate may not find a minimum, but a small learning rate might get stuck in a local minimum

---

# Backpropagation

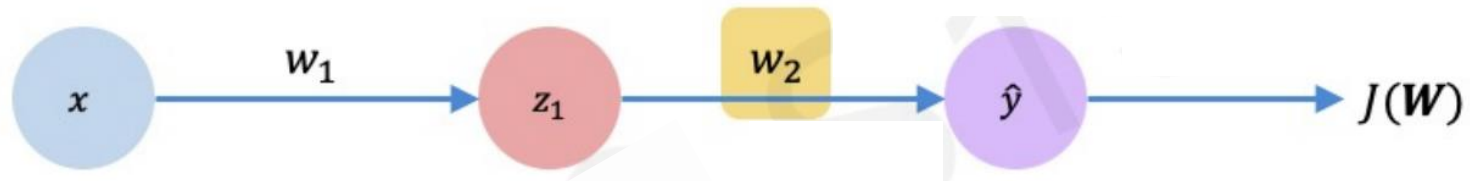


How does a small change in one weight (ex.  $w_2$ ) affect the final loss  $J(\mathbf{w})$ ?



---

# Backpropagation



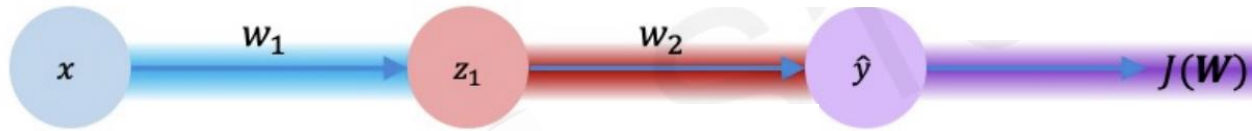
How does a small change in one weight (ex.  $w_2$ ) affect the final loss  $J(\mathbf{W})$ ?

$$\frac{\partial J(\mathbf{W})}{\partial w_2} = \frac{\partial J(\mathbf{W})}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_2}$$

The first term  $\frac{\partial J(\mathbf{W})}{\partial \hat{y}}$  is underlined in purple, and the second term  $\frac{\partial \hat{y}}{\partial w_2}$  is underlined in red.

---

# Backpropagation



$$\frac{\partial J(W)}{\partial w_1} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{purple}} * \underbrace{\frac{\partial \hat{y}}{\partial z_1}}_{\text{red}} * \underbrace{\frac{\partial z_1}{\partial w_1}}_{\text{blue}}$$

Applying the chain rule !

---

# Recommended resources

- Stanford's *Convolutional neural networks for visual recognition* notes on neural networks  
<https://cs231n.github.io/neural-networks-1/>
- MIT's intro to deep learning course  
<https://introtodeeplearning.com/>
- A Tutorial on Deep Learning Part 2: Autoencoders, Convolutional Neural Networks and Recurrent Neural Networks [tutorial2.pdf](#)